Q1: Are we allowed to use Python for the B+ Tree project?

Ans: Any programming language can be used as long as the size information could be collected and presented in a clear way. Note that B+ tree structure uses the concepts of pointers quite often. Projects based on different programming languages will be assessed equally.

Q2: When packing fields into records, is it okay to have a fixed format with fixed length? Will doing so cause marks to be penalized?

Ans: The way how the packing is done is flexible. As long as the design could be well justified, there should be no penalty.

Q3: Is it required to consider other data types than those in the dataset?

Ans: It is good enough to consider only the data types in the dataset.

Q4: Is it required to consider other datasets than the provided dataset?

Ans: It is good enough to consider only the provided dataset.

Q5: The experiments in Project 1, should be executed on separate instances or sequentially (ie: run experiment 1 then experiment 2, etc...).?

Ans: Please do sequentially.

Q6: For the storage requirement, does this mean that multiple nodes can be stored on a single block or must each node be stored on a separate block?

Ans: Each node is stored in main memory for simplicity, but the size occupied by each node should be bounded by that of a block.

Q7: What is meant by a database file and how should it be described?

Ans: A database file means the collection of all the blocks storing a table. There are several options of organizing these blocks – please refer to the lecture slides. Thus, the description of the database file means to specify the options you have designed/adopted.

Q8: Are we going to be assessed on how efficient our storage component is (hence we should minimize the number of bytes our implementation takes) or is any reasonable size acceptable as long as we can explain the reasoning behind it in the report?

Ans: It is the latter case and there is no requirement to do optimization on the storage efficiency.